Package 'iml'

February 5, 2019

Type Package

Title Interpretable Machine Learning

Version 0.9.0

Date 2019-02-01

Maintainer Christoph Molnar <christoph.molnar@gmail.com>

Description Interpretability methods to analyze the behavior and predictions of
any machine learning model.
Implemented methods are:
Feature importance described by Fisher et al. (2018) <arXiv:1801.01489>,
accumulated local effects plots described by Apley (2018) <arXiv:1612.08468>,
partial dependence plots described by Friedman (2001) <http://www.jstor.org/stable/2699986>,
individual conditional expectation ('ice') plots described by Gold-
stein et al. (2013) <doi:10.1080/10618600.2014.907095>,
local models (variant of 'lime') described by Ribeiro et. al (2016) <arXiv:1602.04938>,
the Shapley Value described by Strumbelj et. al (2014) <doi:10.1007/s10115-013-0679-x>,
feature interactions described by Friedman et. al <doi:10.1214/07-AOAS148> and
tree surrogate models.

URL <https://github.com/christophM/iml>

BugReports <https://github.com/christophM/iml/issues>

Imports R6, checkmate, ggplot2, partykit, glmnet, Metrics, data.table,
foreach, yaImpute, prediction, Formula, gridExtra

Suggests randomForest, gower, testthat, rpart, MASS, caret, e1071,
knitr, mlr, covr, rmarkdown, devtools, doParallel, ALEPlot,
ranger

License MIT + file LICENSE

Encoding UTF-8

RoxygenNote 6.1.1

VignetteBuilder knitr

NeedsCompilation no

Author Christoph Molnar [aut, cre]

Repository CRAN

Date/Publication 2019-02-05 18:43:18 UTC

R topics documented:

---

iml-package   *Make machine learning models and predictions interpretable*

Description

The iml package provides tools to analyse machine learning models and predictions.

Author(s)

Maintainer: Christoph Molnar <christoph.molnar@gmail.com>

See Also

[Book on Interpretable Machine Learning](#)

---

calculate.ale.cat	*Compute ALE for 1 categorical feature*

Description

Compute ALE for 1 categorical feature

Usage

```
calculate.ale.cat(dat, run.prediction, feature.name)
```

Arguments

dat	the data.frame with same columns as training data
run.prediction	Predict function of type: f(newdata)
feature.name	The column name of the feature for which to compute ALE

---

calculate.ale.num	*Compute ALE for 1 numerical feature*

Description

Compute ALE for 1 numerical feature

Usage

```
calculate.ale.num(dat, run.prediction, feature.name, grid.size)
```

Arguments

dat	the data.frame with same columns as training data
run.prediction	Predict function of type: f(newdata)
feature.name	The column name of the feature for which to compute ALE
grid.size	The number of intervals

---

calculate.ale.num.cat   *Compute ALE for 2 features, one numerical, one categorical*

Description

Compute ALE for 2 features, one numerical, one categorical

Usage

```
calculate.ale.num.cat(dat, run.prediction, feature.name, grid.size)
```

Arguments

`dat`	the data.frame with same columns as training data
`run.prediction`	Predict function of type: f(newdata)
`feature.name`	The column name of the features for which to compute ALE
`grid.size`	The number of intervals for the numerical feature

---

calculate.ale.num.num   *Compute ALE for 2 numerical features*

Description

Compute ALE for 2 numerical features

Usage

```
calculate.ale.num.num(dat, run.prediction, feature.name, grid.size)
```

Arguments

`dat`	the data.frame with same columns as training data
`run.prediction`	Predict function of type: f(newdata)
`feature.name`	The column names of the feature for which to compute ALE
`grid.size`	The number of cells

FeatureEffect   *Effect of a feature on predictions*

Description

FeatureEffect computes and plots (individual) feature effects of prediction models.

Format

[R6Class](#) object.

Usage

```
effect = FeatureEffect$new(predictor, feature, method = "ale",
    grid.size = 20,  center.at = NULL, run = TRUE)

plot(effect)
effect$results
print(effect)
effectd$set.feature("x2")
```

Arguments

For FeatureEffect$new():

predictor: (Predictor)
> The object (created with Predictor$new()) holding the machine learning model and the data.

feature: ('character(1)' | 'character(2)' | 'numeric(1)' | 'numeric(2)')
> The feature name or index for which to compute the effects.

method: ('character(1)')
> 'ale' for accumulated local effects (the default), 'pdp' for partial dependence plot, 'ice' for individual conditional expectation curves, 'pdp+ice' for partial dependence plot and ice curves within the same plot.

center.at: ('numeric(1)')
> Value at which the plot should be centered. Ignored in the case of two features.

grid.size: ('numeric(1)' | 'numeric(2)')
> The size of the grid for evaluating the predictions

Details

The FeatureEffect class compute the effect a feature has on the prediction. Different methods are implemented:

- Accumulated Local Effect (ALE) plots
- Partial Dependence Plots (PDPs)
- Individual Conditional Expectation (ICE) curves

Accumuluated local effects and partial dependence plots both show the average model prediction over the feature. The difference is that ALE are computed as accumulated differences over the conditional distribution and partial dependence plots over the marginal distribution. ALE plots preferable to PDPs, because they are faster and unbiased when features are correlated.

ALE plots for categorical features are automatically ordered by the similarity of the categories based on the distribution of the other features for instances in a category. When the feature is an ordered factor, the ALE plot leaves the order as is.

Individual conditional expectation curves describe how, for a single observation, the prediction changes when the feature changes and can be combined with partial dependence plots.

To learn more about accumulated local effects, read the Interpretable Machine Learning book: `https://christophm.github.io/interpretable-ml-book/ale.html`

And for the partial dependence plot: `https://christophm.github.io/interpretable-ml-book/pdp.html`

And for individual conditional expectation: `https://christophm.github.io/interpretable-ml-book/ice.html`

Fields

method: ('character(1)')
   'ale' for accumulated local effects, 'pdp' for partial dependence plot, 'ice' for individual conditional expectation curves, 'pdp+ice' for partial dependence plot and ice curves within the same plot.

feature.name: ('character(1)' | 'character(2)')
   The names of the features for which the partial dependence was computed.

feature.type: ('character(1)' | 'character(2)')
   The detected types of the features, either "categorical" or "numerical".

grid.size: ('numeric(1)' | 'numeric(2)')
   The size of the grid.

center.at: ('numeric(1)' | 'character(1)')
   The value for the centering of the plot. Numeric for numeric features, and the level name for factors.

n.features: ('numeric(1)')
   The number of features (either 1 or 2)

predictor: (Predictor)
   The prediction model that was analysed.

results: (data.frame)
   data.frame with the grid of feature of interest and the predicted $\hat{y}$. Can be used for creating custom effect plots.

Methods

center() method to set the value at which the ice computations are centered. See examples.

set.feature() method to get/set feature(s) (by index) fpr which to compute pdp. See examples for usage.

plot() method to plot the effects. See plot.FeatureEffect

predict() method to predict the marginal outcome given a feature. Accepts a data.frame with the feature or a vector. Returns the values of the effect curves at the given values.

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

References

Apley, D. W. 2016. "Visualizing the Effects of Predictor Variables in Black Box Supervised Learning Models." ArXiv Preprint.

Friedman, J.H. 2001. "Greedy Function Approximation: A Gradient Boosting Machine." Annals of Statistics 29: 1189-1232.

Goldstein, A., Kapelner, A., Bleich, J., and Pitkin, E. (2013). Peeking Inside the Black Box: Visualizing Statistical Learning with Plots of Individual Conditional Expectation, 1-22. https://doi.org/10.1080/10618600.2014.907

See Also

plot.FeatureEffect

Examples

```
# We train a random forest on the Boston dataset:
if (require("randomForest")) {
data("Boston", package  = "MASS")
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
mod = Predictor$new(rf, data = Boston)

# Compute the accumulated local effects for the first feature
eff = FeatureEffect$new(mod, feature = "rm",grid.size = 30)
eff$plot()

# Again, but this time with a partial dependence plot and ice curves
eff = FeatureEffect$new(mod, feature = "rm", method = "pdp+ice", grid.size = 30)
plot(eff)

# Since the result is a ggplot object, you can extend it:
if (require("ggplot2")) {
 plot(eff) +
 # Adds a title
 ggtitle("Partial dependence") +
 # Adds original predictions
 geom_point(data = Boston, aes(y = mod$predict(Boston)[[1]], x = rm),
 color =  "pink", size = 0.5)
}

# If you want to do your own thing, just extract the data:
eff.dat = eff$results
head(eff.dat)

# You can also use the object to "predict" the marginal values.
eff$predict(Boston[1:3,])
```

```
# Instead of the entire data.frame, you can also use feature values:
eff$predict(c(5,6,7))

# You can reuse the pdp object for other features:
eff$set.feature("lstat")
plot(eff)

# Only plotting the aggregated partial dependence:
eff = FeatureEffect$new(mod, feature = "crim", method = "pdp")
eff$plot()

# Only plotting the individual conditional expectation:
eff = FeatureEffect$new(mod, feature = "crim", method = "ice")
eff$plot()

# Accumulated local effects and partial dependence plots support up to two features:
eff = FeatureEffect$new(mod, feature = c("crim", "lstat"))
plot(eff)


# FeatureEffect plots also works with multiclass classification
rf = randomForest(Species ~ ., data = iris, ntree=50)
mod = Predictor$new(rf, data = iris, type = "prob")

# For some models we have to specify additional arguments for the predict function
plot(FeatureEffect$new(mod, feature = "Petal.Width"))

# FeatureEffect plots support up to two features:
eff = FeatureEffect$new(mod, feature = c("Sepal.Length", "Petal.Length"))
eff$plot()

# show where the actual data lies
eff$plot(show.data = TRUE)

# For multiclass classification models, you can choose to only show one class:
mod = Predictor$new(rf, data = iris, type = "prob", class = 1)
plot(FeatureEffect$new(mod, feature = "Sepal.Length"))
}
```

FeatureEffects   *Effect of all features on predictions*

Description

FeatureEffects computes feature effects for multiple features at once.

Format

[R6Class](#) object.

Usage

```
effects = FeatureEffects$new(predictor, features = NULL, method = "ale",
    grid.size = 20, center.at = NULL, parallel = FALSE)

plot(effects)
effects$results
print(effects)
```

Arguments

For FeatureEffects$new():

predictor: (Predictor)
> The object (created with Predictor$new()) holding the machine learning model and the data.

features: ('character()')
> The names of the features for which the effects should be computed. Default is all features used in the prediction model.

method: ('character(1)')
> 'ale' for accumulated local effects (the default), 'pdp' for partial dependence plot, 'ice' for individual conditional expectation curves, 'pdp+ice' for partial dependence plot and ice curves within the same plot.

center.at: ('numeric(1)')
> Value at which the plot should be centered.

grid.size: ('numeric(1)' | 'numeric(2)')
> The size of the grid for evaluating the predictions

parallel: 'logical(1)'
> Should the method be executed in parallel? If TRUE, requires a cluster to be registered, see ?foreach::foreach.

Details

FeatureEffects computes the effects for all given features on the model prediction. FeatureEffects is a convenience class that calls FeatureEffect multiple times. See ?FeatureEffect for details what's actually computed.

Only first-order effects can be computed with the FeatureEffects interface. If you are intereated in the visualization of interactions between two features, directly use FeatureEffect.

Fields

method: ('character(1)')
> 'ale' for accumulated local effects, 'pdp' for partial dependence plot, 'ice' for individual conditional expectation curves, 'pdp+ice' for partial dependence plot and ice curves within the same plot.

features: ('character()')
> The names of the features for which the effects were computed.

grid.size: ('numeric(1)' | 'numeric(2)')
> The size of the grid.

center.at: ('numeric(1)' | 'character(1)')
>   The value for the centering of the plot. Numeric for numeric features, and the level name for factors.

predictor: (Predictor)
>   The prediction model that was analysed.

results: (list)
>   A list with the results of each feature effect. Each entry is a data.frame with the grid of feature of interest and the predicted $\hat{y}$. Can be used for creating custom effect plots.

effects: (list)
>   A list of the FeatureEffect objects for each feature. See ?FeatureEffect what you can do with them (e.g. plot them individually).

Methods

plot() method to plot the all effects. See [plot.FeatureEffects](#)

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

References

Apley, D. W. 2016. "Visualizing the Effects of Predictor Variables in Black Box Supervised Learning Models." ArXiv Preprint.

Friedman, J.H. 2001. "Greedy Function Approximation: A Gradient Boosting Machine." Annals of Statistics 29: 1189-1232.

Goldstein, A., Kapelner, A., Bleich, J., and Pitkin, E. (2013). Peeking Inside the Black Box: Visualizing Statistical Learning with Plots of Individual Conditional Expectation, 1-22. https://doi.org/10.1080/10618600.2014.907

See Also

[plot.FeatureEffect](#)

Examples

```
# We train a random forest on the Boston dataset:
if (require("rpart")) {
data("Boston", package  = "MASS")
rf = rpart(medv ~ ., data = Boston)
mod = Predictor$new(rf, data = Boston)

# Compute the accumulated local effects for all features
eff = FeatureEffects$new(mod)
eff$plot()

## Not run:
# Again, but this time with a partial dependence plot
eff = FeatureEffects$new(mod, method = "pdp")
eff$plot()

# Only a subset of features
```

```
eff = FeatureEffects$new(mod, features = c("nox", "crim"))
eff$plot()


# You can access each FeatureEffect individually

eff.nox = eff$effects[["nox"]]
eff.nox$plot()


# FeatureEffects also works with multiclass classification
rf = rpart(Species ~ ., data = iris)
mod = Predictor$new(rf, data = iris, type = "prob")

FeatureEffects$new(mod)$plot(ncol = 2)

## End(Not run)
}
```

FeatureImp   *Feature importance*

Description

FeatureImp computes feature importances for prediction models. The importance is measured as the factor by which the model's prediction error increases when the feature is shuffled.

Format

[R6Class](#) object.

Usage

```
imp = FeatureImp$new(predictor, loss, compare = "ratio", n.repetitions = 5)

imp$plot()
imp$results
print(imp)
```

Arguments

For FeatureImp$new():

predictor: (Predictor)
   The object (created with Predictor$new()) holding the machine learning model and the data.

loss: ('character(1)' | function)
   The loss function. Either the name of a loss (e.g. "ce" for classification or "mse") or a function. See Details for allowed losses.

compare: ('character(1)')

Either "ratio" or "difference". Should importance be measured as the difference or as the ratio of original model error and model error after permutation? Ratio: error.permutation/error.orig, Difference: error.permutation - error.orig

n.repetitions: 'numeric(1)'

How often should the shuffling of the feature be repeated? The higher the number of repetitions the more stable and accurate the results become.

parallel: 'logical(1)'

Should the method be executed in parallel? If TRUE, requires a cluster to be registered, see ?foreach::foreach.

Details

To compute the feature importance for a single feature, the model prediction loss (error) is measured before and after shuffling the values of the feature. By shuffling the feature values, the association between the outcome and the feature is destroyed. The larger the increase in prediction error, the more important the feature was. The shuffling is repeated to get more accurate results, since the permutation feature importance tends to be quite instable. Read the Interpretable Machine Learning book to learn about feature importance in detail: https://christophm.github.io/interpretable-ml-book/feature-importance.html

The loss function can be either specified via a string, or by handing a function to `FeatureImp()`. If you want to use your own loss function it should have this signature: function(actual, predicted). Using the string is a shortcut to using loss functions from the `Metrics` package. Only use functions that return a single performance value, not a vector. Allowed losses are: "ce", "f1", "logLoss", "mae", "mse", "rmse", "mape", "mdae", "msle", "percent_bias", "rae", "rmse", "rmsle", "rse", "rrse", "smape" See `library(help = "Metrics")` to get a list of functions.

Fields

original.error: ('numeric(1)')

The loss of the model before perturbing features.

predictor: (Predictor)

The prediction model that was analysed.

compare: ('character(1)')

Either "ratio" or "difference", depending on whether the importance was calculated as difference between original model error and model error after permutation or as ratio.

results: (data.frame)

data.frame with the results of the feature importance computation. One row per feature with the following columns: importance.05 (5 importance (median importance), importance.95 (95 is also visualized as a bar in the plots, the median importance over the repetitions as a point.

Methods

loss(actual,predicted) The loss function. Can also be applied to data: `object$loss(actual, predicted)`

plot() method to plot the feature importances. See plot.FeatureImp

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

References

Fisher, A., Rudin, C., and Dominici, F. (2018). Model Class Reliance: Variable Importance Measures for any Machine Learning Model Class, from the "Rashomon" Perspective. Retrieved from http://arxiv.org/abs/1801.01489

Examples

```
if (require("rpart")) {
# We train a tree on the Boston dataset:
data("Boston", package  = "MASS")
tree = rpart(medv ~ ., data = Boston)
y = Boston$medv
X = Boston[-which(names(Boston) == "medv")]
mod = Predictor$new(tree, data = X, y = y)


# Compute feature importances as the performance drop in mean absolute error
imp = FeatureImp$new(mod, loss = "mae")

# Plot the results directly
plot(imp)


# Since the result is a ggplot object, you can extend it:
if (require("ggplot2")) {
  plot(imp) + theme_bw()
  # If you want to do your own thing, just extract the data:
  imp.dat = imp$results
  head(imp.dat)
  ggplot(imp.dat, aes(x = feature, y = importance)) + geom_point() +
  theme_bw()
}

# We can also look at the difference in model error instead of the ratio
imp = FeatureImp$new(mod, loss = "mae", compare = "difference")

# Plot the results directly
plot(imp)


# FeatureImp also works with multiclass classification.
# In this case, the importance measurement regards all classes
tree = rpart(Species ~ ., data= iris)
X = iris[-which(names(iris) == "Species")]
y = iris$Species
mod = Predictor$new(tree, data = X, y = y, type = "prob")

# For some models we have to specify additional arguments for the predict function
imp = FeatureImp$new(mod, loss = "ce")
plot(imp)

# For multiclass classification models, you can choose to only compute performance for one class.
```

```
# Make sure to adapt y
mod = Predictor$new(tree, data = X, y = y == "virginica",
 type = "prob", class = "virginica")
imp = FeatureImp$new(mod, loss = "ce")
plot(imp)
}
```

---

impute_cells   *Impute missing cells of grid*

Description

by default assumes first column of cell.dat is x1 and second is x2 leave grid1 NULL if feature x1 is a factor the difference variable has to be named .yhat.diff

Usage

```
impute_cells(cell.dat, grid1 = NULL, grid2, x1.ind = 1, x2.ind = 2)
```

Arguments

cell.dat	data.table with at least 4 columns: .yhat.diff and the two interval indices. Make sure that empty cells are also included and cell.dat is not the sparse representation.
grid1	data.frame where each row is the actual value for a given interval index for feature 1. If empty impute_cells assumes that the feature is categorical (factor).
grid2	data.frame where each row is the actual value for a given interval index for feature 2
x1.ind	column number or name of cell.dat for feature 1. If one feature is categorical, has to be x1
x2.ind	column number or name of cell.dat for feature 2

Interaction   *Feature interactions*

Description

Interaction estimates the feature interactions in a prediction model.

Format

[R6Class](#) object.

Usage

```
ia = Interaction$new(predictor, feature = NULL, grid.size = 20)

plot(ia)
ia$results
print(ia)
```

Arguments

For Interaction$new():

predictor: (Predictor)
   The object (created with Predictor$new()) holding the machine learning model and the data.

feature: ('numeric(1)'|NULL)
   If NULL, for each feature the interactions with all other features are estimated. If one feature name is selected, the 2-way interactions of this feature with all other features are estimated

grid.size: ('logical(1)')
   The number of values per feature that should be used to estimate the interaction strength. A larger grid.size means more accurate the results but longer the computation time. For each of the grid points, the partial dependence functions have to be computed, which involves marginalizing over all data points.

parallel: 'logical(1)'
   Should the method be executed in parallel? If TRUE, requires a cluster to be registered, see ?foreach::foreach.

Details

Interactions between features are measured via the decomposition of the prediction function: If a feature j has no interaction with any other feature, the prediction function can be expressed as the sum of the partial function that depends only on j and the partial function that only depends on features other than j. If the variance of the full function is completely explained by the sum of the partial functions, there is no interaction between feature j and the other features. Any variance that is not explained can be attributed to the interaction and is used as a measure of interaction strength.

The interaction strength between two features is the proportion of the variance of the 2-dimensional partial dependence function that is not explained by the sum of the two 1-dimensional partial dependence functions.

The interaction is measured by Friedman's H-statistic (square root of the H-squared test statistic) and takes on values between 0 (no interaction) to 1 (100

To learn more about interaction effects, read the Interpretable Machine Learning book: https://christophm.github.io/interpretable-ml-book/interaction.html

Fields

grid.size: ('logical(1)')
   The number of values per feature that should be used to estimate the interaction strength.

predictor: (Predictor)
   The prediction model that was analysed.

results: (data.frame)
>   Data.frame with the interaction strength (column '.interation') per feature calculated as Fried-
>   man's H-statistic and - in the case of a multi-dimensional outcome - per class.

Methods

plot() method to plot the feature interactions. See plot.Interaction.

clone() [internal] method to clone the R6 object.

initialize() [internal] method to initialize the R6 object.

References

Friedman, Jerome H., and Bogdan E. Popescu. "Predictive learning via rule ensembles." The Annals
of Applied Statistics 2.3 (2008): 916-954.

Examples

```
## Not run:
if (require("rpart")) {
set.seed(42)
# Fit a CART on the Boston housing data set
data("Boston", package  = "MASS")
rf = rpart(medv ~ ., data = Boston)
# Create a model object
mod = Predictor$new(rf, data = Boston[-which(names(Boston) == "medv")])

# Measure the interaction strength
ia = Interaction$new(mod)

# Plot the resulting leaf nodes
plot(ia)


# Extract the results
dat = ia$results
head(dat)

# Interaction also works with multiclass classification
rf = rpart(Species ~ ., data = iris)
mod = Predictor$new(rf, data = iris, type = "prob")

# For some models we have to specify additional arguments for the predict function
ia = Interaction$new(mod)

ia$plot()

# For multiclass classification models, you can choose to only show one class:
mod = Predictor$new(rf, data = iris, type = "prob", class = "virginica")
plot(Interaction$new(mod))
}

## End(Not run)
```

LocalModel   *LocalModel*

Description

`LocalModel` fits locally weighted linear regression models (logistic regression for classification) to explain single predictions of a prediction model.

Format

[R6Class](#) object.

Usage

```
lime = LocalModel$new(predictor, x.interest = NULL, dist.fun = "gower",
                      kernel.width = NULL, k = 3)

plot(lime)
predict(lime, newdata)
lime$results
lime$explain(x.interest)
print(lime)
```

Arguments

For LocalModel$new():

predictor: (Predictor)
   The object (created with Predictor$new()) holding the machine learning model and the data.

x.interest: (data.frame)
   Single row with the instance to be explained.

dist.fun: ('character(1)')
   The name of the distance function for computing proximities (weights in the linear model). Defaults to "gower". Otherwise will be forwarded to [stats::dist].

kernel.width: ('numeric(1)')
   The width of the kernel for the proximity computation. Only used if dist.fun is not 'gower'.

k: ('numeric(1)')
   The (maximum) number of features to be used for the surrogate model.

Details

A weighted glm is fitted with the machine learning model prediction as target. Data points are weighted by their proximity to the instance to be explained, using the gower proximity measure. L1-regularisation is used to make the results sparse. The resulting model can be seen as a surrogate for the machine learning model, which is only valid for that one point. Categorical features are binarized, depending on the category of the instance to be explained: 1 if the category is the same,

0 otherwise. To learn more about local models, read the Interpretable Machine Learning book: https://christophm.github.io/interpretable-ml-book/lime.html

The approach is similar to LIME, but has the following differences:

- Distance measure: Uses as default the gower proximity (= 1 - gower distance) instead of a kernel based on the Euclidean distance. Has the advantage to have a meaningful neighbourhood and no kernel width to tune. When the distance is not "gower", then the stats::dist() function with the chosen method will be used, and turned into a similarity measure: sqrt(exp(-(distance^2) / (kernel.width^2))).
- Sampling: Uses the original data instead of sampling from normal distributions. Has the advantage to follow the original data distribution.
- Visualisation: Plots effects instead of betas. Both are the same for binary features, but ared different for numerical features. For numerical features, plotting the betas makes no sense, because a negative beta might still increase the prediction when the feature value is also negative.

To learn more about local surrogate models, read the Interpretable Machine Learning book: https://christophm.github.io/interpretable-ml-book/lime.html

Fields

best.fit.index: ('numeric(1)')
> The index of the best glmnet fit.

k: ('numeric(1)')
> The number of features as set by the user.

model: (glmnet)
> The fitted local model.

predictor: (Predictor)
> The prediction model that was analysed.

results: (data.frame)
> Results with the feature names (`feature`) and contributions to the prediction

x.interest: (data.frame)
> The instance to be explained. See Examples for usage.

Methods

explain(x.interest) method to set a new data point which to explain.

plot() method to plot the LocalModel feature effects. See plot.LocalModel

predict() method to predict new data with the local model See also predict.LocalModel

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

References

Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why Should I Trust You?": Explaining the Predictions of Any Classifier. Retrieved from http://arxiv.org/abs/1602.04938

Gower, J. C. (1971), "A general coefficient of similarity and some of its properties". Biometrics, 27, 623–637.

See Also

[plot.LocalModel](#) and [predict.LocalModel](#)

[Shapley](#) can also be used to explain single predictions

The package 'lime' with the original implementation

Examples

```
if (require("randomForest")) {
# First we fit a machine learning model on the Boston housing data
data("Boston", package  = "MASS")
X = Boston[-which(names(Boston) == "medv")]
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
mod = Predictor$new(rf, data = X)

# Explain the first instance of the dataset with the LocalModel method:
x.interest = X[1,]
lemon = LocalModel$new(mod, x.interest = x.interest, k = 2)
lemon

# Look at the results in a table
lemon$results
# Or as a plot
plot(lemon)

# Reuse the object with a new instance to explain
lemon$x.interest
lemon$explain(X[2,])
lemon$x.interest
plot(lemon)

# LocalModel also works with multiclass classification
rf = randomForest(Species ~ ., data= iris, ntree=50)
X = iris[-which(names(iris) == 'Species')]
mod = Predictor$new(rf, data = X, type = "prob", class = "setosa")

# Then we explain the first instance of the dataset with the LocalModel method:
lemon = LocalModel$new(mod, x.interest = X[1,], k = 2)
lemon$results
plot(lemon)
}
```

---

order_levels   *Order levels of a categorical features*

Description

Goal: Compute the distances between two categories. Input: Instances from category 1 and 2

Usage

```
order_levels(dat, feature.name)
```

Arguments

`dat`	data.frame with the training data
`feature.name`	the name of the categorical feature

Details

1. For all features, do (excluding the categorical feature for which we are computing the order): - If the feature is numerical: Take instances from category 1, calculate the empirical cumulative probability distribution function (ecdf) of the feature. The ecdf is a function that tells us for a given feature value, how many values are smaller. Do the same for category 2. The distance is the absolute maximum point-wise distance of the two ecdf. Practically, this value is high when the distribution from one category is strongly shifted far away from the other. This measure is also known as the [Kolmogorov-Smirnov distance](https://en.wikipedia.org/wiki/Kolmogorov - If the feature is categorical: Take instances from category 1 and calculate a table with the relative frequency of each category of the other feature. Do the same for instances from category 2. The distance is the sum of the absolute difference of both relative frequency tables. 2. Sum up the distances over all features

This algorithm we run for all pairs of categories. Then we have a k times k matrix, when k is the number of categories, where each entry is the distance between two categories. Still not enough to have a single order, because, a (dis)similarity tells you the pair-wise distances, but does not give you a one-dimensional ordering of the classes. To kind of force this thing into a single dimension, we have to use a dimension reduction trick called multi-dimensional scaling. This can be solved using multi-dimensional scaling, which takes in a distance matrix and returns a distance matrix with reduced dimension. In our case, we only want 1 dimension left, so that we have a single ordering of the categories and can compute the accumulated local effects. After reducing it to a single ordering, we are done and can use this ordering to compute ALE. This is not the Holy Grail how to order the factors, but one possibility.

Orders the levels by their similarity in other features. Computes per feature the distance, sums up all distances and does multi-dimensional scaling

Value

the order of the levels (not levels itself)

Partial   *Effect of one or two feature(s) on the model predictions (deprecated)*

Description

Deprecated, please use 'FeatureEffect', see ?FeatureEffect

Format

R6Class object.

See Also

[FeatureEffect](#)

```
plot.FeatureEffect        Plot FeatureEffect
```

Description

plot.FeatureEffect() plots the results of a FeatureEffect object.

Usage

```
## S3 method for class 'FeatureEffect'
plot(x, rug = TRUE, show.data = FALSE,
  ylim = NULL)
```

Arguments

x	A FeatureEffect R6 object
rug	[logical] Should a rug be plotted to indicate the feature distribution? The rug will be jittered a bit, so the location may not be exact, but it avoids overplotting.
show.data	Should the data points be shown? Only affects 2D plots, and ignored for 1D plots, because rug has the same information.
ylim	Vector with two coordinates for the y-axis. Only works when one feature is used in FeatureEffet, ignored when two are used.

Value

ggplot2 plot object

See Also

[FeatureEffect](#)

Examples

```
# We train a random forest on the Boston dataset:
if (require("randomForest")) {
data("Boston", package  = "MASS")
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
mod = Predictor$new(rf, data = Boston)

# Compute the ALE for the first feature
eff = FeatureEffect$new(mod, feature = "crim")

# Plot the results directly
plot(eff)
}
```

---

plot.FeatureEffects   *Plot FeatureEffect*

Description

plot.FeatureEffect() plots the results of a FeatureEffect object.

Usage

```
## S3 method for class 'FeatureEffects'
plot(x, features = NULL, nrows = NULL,
  ncols = NULL, fixed_y = TRUE, ...)
```

Arguments

x	A FeatureEffect R6 object
features	[character()] For which features should the effects be plotted? Default is all features. You can also sort the order of the plots with this argument.
nrows	The number of rows in the table of graphics
ncols	The number of columns in the table of graphics
fixed_y	Should the y-axis range be the same for all effects? Defaults to TRUE.
...	Further arguments for FeatureEffect$plot()

Details

In contrast to other plot methods in iml, for FeatureEffects the returned plot is not a ggplot2 object, but a grid object, a collection of multiple ggplot2 plots.

Value

grid object

See Also

[FeatureEffects plot.FeatureEffect](#)

Examples

```
# We train a random forest on the Boston dataset:
if (require("randomForest")) {
data("Boston", package = "MASS")
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
mod = Predictor$new(rf, data = Boston)

# Compute the partial dependence for the first feature
eff = FeatureEffects$new(mod)
```

```
# Plot the results directly
eff$plot()

# For a subset of features
eff$plot(features = c("lstat", "crim"))

# With a different layout
eff$plot(nrows = 2)
}
```

---

plot.FeatureImp   *Plot Feature Importance*

Description

plot.FeatureImp() plots the feature importance results of a FeatureImp object.

Usage

```
## S3 method for class 'FeatureImp'
plot(x, sort = TRUE, ...)
```

Arguments

x	A FeatureImp R6 object
sort	logical. Should the features be sorted in descending order? Defaults to TRUE.
...	Further arguments for the objects plot function

Details

The plot shows the importance per feature.

When n.repetitions in 'FeatureImp$new' was larger than 1, then we get multiple importance esti-mates per feature. The importances are aggregated and the plot shows the median importance per feature (as dots) and also the 90 helps to understand how much variance the computation has per feature.

Value

ggplot2 plot object

See Also

[FeatureImp](#)

Examples

```
if (require("rpart")) {
# We train a tree on the Boston dataset:
data("Boston", package  = "MASS")
tree = rpart(medv ~ ., data = Boston)
y = Boston$medv
X = Boston[-which(names(Boston) == "medv")]
mod = Predictor$new(tree, data = X, y = y)

# Compute feature importances as the performance drop in mean absolute error
imp = FeatureImp$new(mod, loss = "mae")

# Plot the results directly
plot(imp)
}
```

---

plot.Interaction   *Plot Interaction*

Description

plot.Interaction() plots the results of an Interaction object.

Usage

```
## S3 method for class 'Interaction'
plot(x, sort = TRUE)
```

Arguments

x	An Interaction R6 object
sort	logical. Should the features be sorted in descending order? Defaults to TRUE.

Value

ggplot2 plot object

See Also

[Interaction](#)

Examples

```
# We train a tree on the Boston dataset:
## Not run:
if (require("rpart")) {
data("Boston", package  = "MASS")
rf = rpart(medv ~ ., data = Boston)
mod = Predictor$new(rf, data = Boston)

# Compute the interactions
ia = Interaction$new(mod)

# Plot the results directly
plot(ia)
}

## End(Not run)
```

---

plot.LocalModel *Plot Local Model*

Description

plot.LocalModel() plots the feature effects of a LocalModel object.

Usage

```
## S3 method for class 'LocalModel'
plot(object)
```

Arguments

object   A LocalModel R6 object

Value

ggplot2 plot object

See Also

[LocalModel](#)

Examples

```
if (require("randomForest")) {
# First we fit a machine learning model on the Boston housing data
data("Boston", package  = "MASS")
X = Boston[-which(names(Boston) == "medv")]
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
mod = Predictor$new(rf, data = X)
```

```
# Explain the first instance of the dataset with the LocalModel method:
x.interest = X[1,]
lemon = LocalModel$new(mod, x.interest = x.interest, k = 2)
plot(lemon)
}
```

---

plot.Shapley	*Plot Shapley*

Description

plot.Shapley() plots the Shapley values - the contributions of feature values to the prediction.

Usage

```
## S3 method for class 'Shapley'
plot(object, sort = TRUE)
```

Arguments

object	A Shapley R6 object
sort	logical. Should the feature values be sorted by Shapley value? Ignored for multi.class output.

Value

ggplot2 plot object

See Also

[Shapley](#)

Examples

```
## Not run:
if (require("rpart")) {
# First we fit a machine learning model on the Boston housing data
data("Boston", package  = "MASS")
rf =  rpart(medv ~ ., data = Boston)
X = Boston[-which(names(Boston) == "medv")]
mod = Predictor$new(rf, data = X)

# Then we explain the first instance of the dataset with the Shapley method:
x.interest = X[1,]
shapley = Shapley$new(mod, x.interest = x.interest)
plot(shapley)
}

## End(Not run)
```

---

plot.TreeSurrogate   *Plot Tree Surrogate*

Description

Plot the response for newdata of a TreeSurrogate object. Each plot facet is one leaf node and visualises the distribution of the $\hat{y}$ from the machine learning model.

Usage

```
## S3 method for class 'TreeSurrogate'
plot(object)
```

Arguments

object   A TreeSurrogate R6 object

Value

ggplot2 plot object

See Also

[TreeSurrogate](#)

Examples

```
if (require("randomForest")) {
# Fit a Random Forest on the Boston housing data set
data("Boston", package  = "MASS")
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
# Create a model object
mod = Predictor$new(rf, data = Boston[-which(names(Boston) == "medv")])

# Fit a decision tree as a surrogate for the whole random forest
dt = TreeSurrogate$new(mod)

# Plot the resulting leaf nodes
plot(dt)
}
```

predict.LocalModel   *Predict LocalModel*

Description

Predict the response for newdata with the LocalModel model.

Usage

```
## S3 method for class 'LocalModel'
predict(object, newdata = NULL, ...)
```

Arguments

object	A LocalModel R6 object
newdata	A data.frame for which to predict
...	Further arguments for the objects predict function

Value

A data.frame with the predicted outcome.

See Also

[LocalModel](#)

Examples

```
if (require("randomForest")) {
# First we fit a machine learning model on the Boston housing data
data("Boston", package  = "MASS")
X = Boston[-which(names(Boston) == "medv")]
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
mod = Predictor$new(rf, data = X)

# Explain the first instance of the dataset with the LocalModel method:
x.interest = X[1,]
lemon = LocalModel$new(mod, x.interest = x.interest, k = 2)
predict(lemon, newdata = x.interest)
}
```

---

predict.TreeSurrogate   *Predict Tree Surrogate*

Description

Predict the response for newdata of a TreeSurrogate object.

Usage

```
## S3 method for class 'TreeSurrogate'
predict(object, newdata, type = "prob", ...)
```

Arguments

object	The surrogate tree. A TreeSurrogate R6 object
newdata	A data.frame for which to predict
type	Either "prob" or "class". Ignored if the surrogate tree does regression.
...	Further argumets for predict_party

Details

This function makes the TreeSurrogate object call its iternal object$predict() method.

Value

A data.frame with the predicted outcome. In case of regression it is the predicted $\hat{y}$. In case of classification it is either the class probabilities (for type "prob") or the class label (type "class")

See Also

[TreeSurrogate](#)

Examples

```
if (require("randomForest")) {
# Fit a Random Forest on the Boston housing data set
data("Boston", package = "MASS")
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
# Create a model object
mod = Predictor$new(rf, data = Boston[-which(names(Boston) == "medv")])

# Fit a decision tree as a surrogate for the whole random forest
dt = TreeSurrogate$new(mod)

# Plot the resulting leaf nodes
predict(dt, newdata = Boston)
}
```

Predictor   *Predictor object*

Description

A `Predictor` object holds any machine learning model (mlr, caret, randomForest, ...) and the data
to be used of analysing the model. The interpretation methods in the iml package need the machine
learning model to be wrapped in a `Predictor` object.

Format

[R6Class](#) object.

Usage

```
model = Predictor$new(model = NULL, data, y = NULL, class=NULL,
  predict.fun = NULL, type = NULL)
```

```
model$predict(newdata)
```

Arguments

model: (any)
   The machine learning model. Recommended are models from mlr and caret. Other machine
   learning with a S3 predict functions work as well, but less robust (e.g. randomForest).

data: (data.frame)
   The data to be used for analysing the prediction model. Allowed column classes are: numeric,
   factor, integer, ordered and character. For some models the data can be extracted automati-
   cally. Predictor$new throws an error when it can't extract the data automatically.

y: (('character(1)') | numeric | factor)
   The target vector or (preferably) the name of the target column in the `data` argument. Predictor
   tries to infer the target automatically from the model.

class: ('character(1)')
   The class column to be returned in case of multiclass output. You can either use numbers, e.g.
   class=2 would take the 2nd column from the predictions, or the column name of the predicted
   class, e.g. class="dog".

predict.fun: (function)
   The function to predict newdata. Only needed if `model` is not a model from mlr or caret
   package. The first argument of predict.fun has to be the model, the second the newdata:
   function(model, newdata)

type: ('character(1)')
   This argument is passed to the prediction function of the model. For regression models you
   usually don't have to provide the type argument. The classic use case is to say type="prob" for
   classification models. Consult the documentation of the machine learning package you use to
   find which type options you have. If both predict.fun and type are used, then type is passed as
   an argument to predict.fun.

batch.size: ('numeric(1)')
The maximum number of rows to be input the model for prediction at once. Currently only respected for FeatureImp, Partial and Interaction.

Details

A Predictor object is a container for the prediction model and the data. This ensures that the machine learning model can be analysed robustly.

Note: In case of classification, the model should return one column per class with the class probability.

Fields

class: ('character(1)')
The class column to be returned.

data: (data.frame)
data object with the data for the model interpretation.

prediction.colnames: (character)
The column names of the predictions.

task: ('character(1)')
The inferred prediction task: "classification" or "regression".

Methods

predict(newdata) method to predict new data with the machine learning model.

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

Examples

```
if (require("mlr")) {
task = makeClassifTask(data = iris, target = "Species")
learner = makeLearner("classif.rpart", minsplit = 7, predict.type = "prob")
mod.mlr = train(learner, task)
mod = Predictor$new(mod.mlr, data = iris)
mod$predict(iris[1:5,])

mod = Predictor$new(mod.mlr, data = iris, class = "setosa")
mod$predict(iris[1:5,])
}

if (require("randomForest")) {
rf = randomForest(Species ~ ., data = iris, ntree = 20)


mod = Predictor$new(rf, data = iris, type = "prob")
mod$predict(iris[50:55,])

# Feature importance needs the target vector, which needs to be supplied:
```

```
mod = Predictor$new(rf, data = iris, y = "Species", type = "prob")
}
```

Shapley   *Prediction explanations with game theory*

Description

`Shapley` computes feature contributions for single predictions with the Shapley value, an approach from cooperative game theory. The features values of an instance cooperate to achieve the prediction. The Shapley value fairly distributes the difference of the instance's prediction and the datasets average prediction among the features.

Format

[R6Class](#) object.

Usage

```
shapley = Shapley$new(predictor, x.interest = NULL, sample.size = 100)

plot(shapley)
shapley$results
print(shapley)
shapley$explain(x.interest)
```

Arguments

For Shapley$new():

predictor: (Predictor)
  The object (created with Predictor$new()) holding the machine learning model and the data.

x.interest: (data.frame)
  Single row with the instance to be explained.

sample.size: ('numeric(1)')
  The number of Monte Carlo samples for estimating the Shapley value.

Details

For more details on the algorithm see https://christophm.github.io/interpretable-ml-book/shapley.html

Fields

predictor: (Predictor)
    The object (created with Predictor$new()) holding the machine learning model and the data.

results: (data.frame)
    data.frame with the Shapley values (phi) per feature.

sample.size: ('numeric(1)')
    The number of times coalitions/marginals are sampled from data X. The higher the more accurate the explanations become.

x.interest: (data.frame)
    Single row with the instance to be explained.

y.hat.interest: (numeric)
    Predicted value for instance of interest

y.hat.average: ('numeric(1)')
    Average predicted value for data X

Methods

explain(x.interest) method to set a new data point which to explain.

plot() method to plot the Shapley value. See plot.Shapley

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

References

Strumbelj, E., Kononenko, I. (2014). Explaining prediction models and individual predictions with feature contributions. Knowledge and Information Systems, 41(3), 647-665. https://doi.org/10.1007/s10115-013-0679-x

See Also

Shapley

A different way to explain predictions: LocalModel

Examples

```
if (require("rpart")) {
# First we fit a machine learning model on the Boston housing data
data("Boston", package  = "MASS")
rf =  rpart(medv ~ ., data = Boston)
X = Boston[-which(names(Boston) == "medv")]
mod = Predictor$new(rf, data = X)

# Then we explain the first instance of the dataset with the Shapley method:
x.interest = X[1,]
shapley = Shapley$new(mod, x.interest = x.interest)
shapley
```

```
# Look at the results in a table
shapley$results
# Or as a plot
plot(shapley)

# Explain another instance
shapley$explain(X[2,])
plot(shapley)
## Not run:
# Shapley() also works with multiclass classification
rf = rpart(Species ~ ., data = iris)
X = iris[-which(names(iris) == "Species")]
mod = Predictor$new(rf, data = X, type = "prob")

# Then we explain the first instance of the dataset with the Shapley() method:
shapley = Shapley$new(mod, x.interest = X[1,])
shapley$results
plot(shapley)

# You can also focus on one class
mod = Predictor$new(rf, data = X, type = "prob", class = "setosa")
shapley = Shapley$new(mod, x.interest = X[1,])
shapley$results
plot(shapley)

## End(Not run)
}
```

TreeSurrogate   *Decision tree surrogate model*

Description

TreeSurrogate fits a decision tree on the predictions of a prediction model.

Format

R6Class object.

Usage

```
tree = TreeSurrogate$new(predictor, maxdepth = 2, tree.args = NULL)

plot(tree)
predict(tree, newdata)
tree$results
print(tree)
```

Arguments

For TreeSurrogate$new():

predictor: (Predictor)
   The object (created with Predictor$new()) holding the machine learning model and the data.

maxdepth: ('numeric(1)')
   The maximum depth of the tree. Default is 2.

tree.args: (named list)
   Further arguments for `ctree`.

Details

A conditional inference tree is fitted on the predicted $\hat{y}$ from the machine learning model and the data. The `partykit` package and function are used to fit the tree. By default a tree of maximum depth of 2 is fitted to improve interpretability.

To learn more about global surrogate models, read the Interpretable Machine Learning book: `https://christophm.github.io/interpretable-ml-book/global.html`

Fields

maxdepth: ('numeric(1)')
   The maximum tree depth.

predictor: (Predictor)
   The prediction model that was analysed.

r.squared: ('numeric(1|n.classes)')
   R squared measures how well the decision tree approximates the underlying model. It is calculated as 1 - (variance of prediction differences / variance of black box model predictions). For the multi-class case, r.squared contains one measure per class.

results: (data.frame)
   Data.frame with sampled feature X together with the leaf node information (columns .node and .path) and the predicted $\hat{y}$ for tree and machine learning model (columns starting with .y.hat).

tree: (party)
   The fitted tree. See also ctree.

Methods

plot() method to plot the leaf nodes of the surrogate decision tree. See plot.TreeSurrogate.

predict() method to predict new data with the tree. See also predict.TreeSurrogate

`clone()` [internal] method to clone the R6 object.

`initialize()` [internal] method to initialize the R6 object.

References

Craven, M., & Shavlik, J. W. (1996). Extracting tree-structured representations of trained networks. In Advances in neural information processing systems (pp. 24-30).

See Also

predict.TreeSurrogate plot.TreeSurrogate

For the tree implementation ctree

Examples

```
if (require("randomForest")) {
# Fit a Random Forest on the Boston housing data set
data("Boston", package = "MASS")
rf = randomForest(medv ~ ., data = Boston, ntree = 50)
# Create a model object
mod = Predictor$new(rf, data = Boston[-which(names(Boston) == "medv")])

# Fit a decision tree as a surrogate for the whole random forest
dt = TreeSurrogate$new(mod)

# Plot the resulting leaf nodes
plot(dt)

# Use the tree to predict new data
predict(dt, Boston[1:10,])

# Extract the results
dat = dt$results
head(dat)


# It also works for classification
rf = randomForest(Species ~ ., data = iris, ntree = 50)
X = iris[-which(names(iris) == "Species")]
mod = Predictor$new(rf, data = X, type = "prob")

# Fit a decision tree as a surrogate for the whole random forest
dt = TreeSurrogate$new(mod, maxdepth=2)

# Plot the resulting leaf nodes
plot(dt)

# If you want to visualise the tree directly:
plot(dt$tree)

# Use the tree to predict new data
set.seed(42)
iris.sample = X[sample(1:nrow(X), 10),]
predict(dt, iris.sample)
predict(dt, iris.sample, type = "class")

# Extract the dataset
dat = dt$results
head(dat)
}
```

Index